

market-segmentation

11, 2023

```
[ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: # Import all required Libraries:  
  
import pandas as pd  
import matplotlib.pyplot as plt  
import re  
import time  
import warnings  
import numpy as np  
from nltk.corpus import stopwords  
from sklearn.decomposition import TruncatedSVD  
from sklearn.preprocessing import normalize  
from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.manifold import TSNE  
import seaborn as sns  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.metrics import confusion_matrix  
from sklearn.metrics import accuracy_score, log_loss  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn.linear_model import SGDClassifier  
from imblearn.over_sampling import SMOTE  
from collections import Counter  
from scipy.sparse import hstack  
from sklearn.multiclass import OneVsRestClassifier  
from sklearn.svm import SVC  
from sklearn.model_selection import StratifiedKFold  
from collections import Counter, defaultdict  
from sklearn.calibration import CalibratedClassifierCV  
from sklearn.naive_bayes import MultinomialNB  
from sklearn.naive_bayes import GaussianNB  
from sklearn.model_selection import train_test_split  
from sklearn.model_selection import GridSearchCV  
import math  
from sklearn.metrics import normalized_mutual_info_score
```

```

from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")
import six
import sys
sys.modules['sklearn.externals.six'] = six
from mlxtend.classifier import StackingClassifier
from sklearn import model_selection
from sklearn.linear_model import LogisticRegression

```

```
[ ]: df = pd.read_csv("/content/drive/MyDrive/out.csv")
```

```
[ ]: df.head()
```

```
[ ]:
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	¥
0	0	1	0	1	0	1	1	0	1	
1	1	1	0	1	1	1	1	1	1	
2	0	1	1	1	1	1	0	1	1	
3	1	1	0	1	1	1	1	1	0	
4	0	1	0	1	1	1	1	0	0	

	healthy	disgusting	Like	Age	VisitFrequency	Gender	Cluster
0	0	0	-3	61		2	0
1	0	0	2	51		2	0
2	1	0	1	62		2	0
3	0	1	4	69		4	0
4	1	0	2	49		3	1

```
[ ]: df.shape
```

```
[ ]: (1453, 16)
```

```
[ ]: df.columns
```

```
[ ]: Index(['yummy', 'convenient', 'spicy', 'fattening', 'greasy', 'fast', 'cheap',
          'tasty', 'expensive', 'healthy', 'disgusting', 'Like', 'Age',
          'VisitFrequency', 'Gender', 'Cluster'],
          dtype='object')
```

```
[ ]: # columns to keep:
data= df[['yummy', 'convenient', 'spicy', 'fattening', 'greasy', 'fast',
          'cheap', 'tasty', 'expensive', 'healthy', 'disgusting', 'Age', 'Gender',
          'Cluster']].rename({'Cluster': 'label'}, axis=1)
```

```
[ ]: data.head(2)
```

```
[ ]:
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	expensive	¥
0	0	1	0	1	0	1	1	0	1	
1	1	1	0	1	1	1	1	1	1	

	healthy	disgusting	Age	Gender	label
0	0	0	61	0	3
1	0	0	51	0	0

```
[ ]: X = data.iloc[:, data.columns != 'label'] y
= data.iloc[:, data.columns == 'label']
```

0.1 Train, Test and Cross-Validation Dataset Construction

```
[ ]: # split the data into test and train by maintaining same distribution of output_
variable 'y_true' [stratify=y_true]
X_train, test_df, y_train, y_test = train_test_split(X, y, stratify=y, _
test_size=0.2)
# split the train data into train and cross validation by maintaining same_
distribution of output variable 'y_train' [stratify=y_train]
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train, _

[ ]: print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0]) print('Number
of data points in cross validation data:', cv_df.shape[0])
```

Number of data points in train data: 929

Number of data points in test data: 291

Number of data points in cross validation data: 233

```
[ ]: test_df.head(2)
```

```
[ ]:
```

	yummy	convenient	spicy	fattening	greasy	fast	cheap	tasty	¥
140	0	1	0	1	0	1	1	1	
1349	1	1	0	1	0	1	0	1	

	expensive	healthy	disgusting	Age	Gender
140	0	0	0	42	1
1349	1	1	0	31	1

```
[ ]: y_test.head(2)
```

```
[ ]:
```

	label
140	2
1349	0

Distribution of y_i's in Train, Test and Cross Validation datasets

```
[ ]: # it returns a dict, keys as class labels and values as the number of data_
    ↪points in that class
train_class_distribution = y_train['label'].value_counts().sort_index()
test_class_distribution = y_test['label'].value_counts().sort_index()
cv_class_distribution = y_cv ['label'].value_counts().sort_index()

my_colors = 'rgbkymc'
train_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in train data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.

# ↪argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing_
    ↪order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', train_class_distribution.
        ↪values[i], '(', np.round((train_class_distribution.values[i]/train_df.
            ↪shape[0]*100), 3), '%')

print('-'*80)
my_colors = 'rgbkymc'
test_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in test data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.

# ↪argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing_
    ↪order
sorted_yi = np.argsort(-test_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', test_class_distribution.
        ↪values[i], '(', np.round((test_class_distribution.values[i]/test_df.
            ↪shape[0]*100), 3), '%')


print('-'*80)
my_colors = 'rgbkymc'
```

```

cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in cross validation data')
plt.grid()
plt.show()


# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
#  $-(\text{train\_class\_distribution.values})$ : the minus sign will give us indecreasing_
# order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', cv_class_distribution.
          values[i], '(', np.round((cv_class_distribution.values[i]/cv_df.
          shape[0]*100), 3), '%')

```



Number of data points in class 3 : 368 (39.612%)
 Number of data points in class 4 : 260 (27.987%)
 Number of data points in class 1 : 193 (20.775)%

Number of data points in class 2 : 108 (11.625 %)




Number of data points in class 3 : 116 (39.863 %)

Number of data points in class 4 : 81 (27.835 %)

Number of data points in class 1 : 60 (20.619 %)

Number of data points in class 2 : 34 (11.684 %)



Number of data points in class 3 : 93 (39.914%)
 Number of data points in class 4 : 65 (27.897%)
 Number of data points in class 1 : 48 (20.601%)
 Number of data points in class 2 : 27 (11.588)%

```
[ ]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
train_df= scaler.fit_transform(train_df)
train_df = pd.DataFrame(train_df) test_df =
= scaler.transform(test_df) test_df =
pd.DataFrame(test_df)
cv_df = scaler.transform(cv_df) cv_df
= pd.DataFrame(cv_df)
```

Prediction using a 'Random' Model

```
[ ]: # This function plots the confusion matrices given  $y_i$ ,  $y_i_{hat}$ .
def plot_confusion_matrix(test_y, predict_y): C
    = confusion_matrix(test_y, predict_y)
```

```

# C = 9, 9 matrix, each cell (i, j) represents number of points of class i
→ are predicted class j
A = ((C.T) / (C.sum(axis=1))).T
#divid each element of the confusion matrix with the sum of elements in
→ that column
# C = [[1, 2],
#      [3, 4]]
# C.T = [[1, 3],
#        [2, 4]]
# C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to
→ rows in two dimensional array #
C.sum(axis = 1) = [[3, 7]]
# ((C.T) / (C.sum(axis=1))) = [[1/3, 3/7] #
                             [2/3, 4/7]]

# ((C.T) / (C.sum(axis=1))).T = [[1/3, 2/3] #
                                [3/7, 4/7]]

# sum of row elements = 1

B = (C / C.sum(axis=0))
#divid each element of the confusion matrix with the sum of elements in
→ that row
# C = [[1, 2],
#      [3, 4]]
# C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds to
→ rows in two dimensional array #
C.sum(axis = 0) = [[4, 6]]
# (C / C.sum(axis=0)) = [[1/4, 2/6],
#                       [3/4, 4/6]]

labels = [1, 2, 3, 4]
# representing A in heatmap format print("-
"*20, "Confusion matrix", "-"*20)
plt.figure(figsize=(20, 7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
→ yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
plt.figure(figsize=(20, 7))
sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
→ yticklabels=labels)

```



```

plt.xlabel('Predicted Class')
plt.ylabel('Original Class') plt.show()

# representing B in heatmap format
print("--*20, \"Recall matrix (Row sum=1)\", \"--*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class') plt.show()

```

[]: *# we need to generate 5 numbers and the sum of numbers should be 1*
one solution is to generate 5 numbers and divide each of the numbers by their

```

sum
# ref: https://stackoverflow.com/a/18662466/4084039
test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,4))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,4)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Cross Validation Data using Random
Model", log_loss(y_cv, cv_predicted_y, eps=1e-15))


# Test-Set error.
#we create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,4))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,4)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random
Model", log_loss(y_test, test_predicted_y, eps=1e-15))
predicted_y = np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y+1)

```


Log loss on Cross Validation Data using Random Model 1.6801866526522178

Log loss on Test Data using Random Model 1.611141956382239


..... Confusion matrix



Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)



Machine Learning Models

[]: *#Data preparation for ML models.*

#Misc. functions for ML models

```
def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we will provide the array of probabilities
    # belongs to each class
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y -
    test_y)) / test_y.shape[0])
    plot_confusion_matrix(test_y, pred_y)
```

[]: **import pickle**
 pickle.dump(sig_clf, open('/content/drive/MyDrive/final_prediction.pickle', 'wb'))
 pickle.dump(scaler, open('/content/drive/MyDrive/scaler.pickle', 'wb'))
def report_log_loss(train_x, train_y, test_x, test_y, clf):
 clf.fit(train_x, train_y)
 sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
 sig_clf.fit(train_x, train_y)
 sig_clf_probs = sig_clf.predict_proba(test_x)
return log_loss(test_y, sig_clf_probs, eps=1e-15)

```
File "<ipython-input-48-bf421f3cc827>", line 3
    pickle.dump(scaler, open('/content/drive/MyDrive/scaler.pickle', 'wb'))def _
report_log_loss(train_x, train_y, test_x, test_y, clf):
```

SyntaxError: invalid syntax

K Nearest Neighbour Classification

Hyper parameter tuning

```
[ ]: alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_df, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_df, y_train)
    sig_clf_probs = sig_clf.predict_proba(cv_df)
    cv_log_error_array.append(log_loss(y_cv, sig_clf_probs, labels=clf.
classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use _
log-probability estimates
    print("Log Loss :", log_loss(y_cv, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_df, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df, y_train)

predict_y = sig_clf.predict_proba(train_df)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
```

```


predict_y = sig_clf.predict_proba(cv_df)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation_
↪ log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_df)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
↪ ", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 5
Log Loss : 0.1903307935393392
for alpha = 11
Log Loss : 0.1777915041878316
for alpha = 15
Log Loss : 0.16426317498928972
for alpha = 21
Log Loss : 0.16391829819291193
for alpha = 31
Log Loss : 0.1371292146305823
for alpha = 41
Log Loss : 0.13658723817517318
for alpha = 51
Log Loss : 0.13771713580884173
for alpha = 99
Log Loss : 0.1575477268166864

```




For values of best alpha = 41 The train log loss is: 0.17882548673489387
 For values of best alpha = 41 The cross validation log loss is:
 0.13658723817517318
 For values of best alpha = 41 The test log loss is: 0.16799609578880664


```
[ ]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
      predict_and_plot_confusion_matrix(train_df.values, y_train.values, cv_df.
      values, y_cv.values, clf)
```


Log loss : 0.13658723817517318
 Number of mis-classified points : 163.81115879828326

Confusion matrix



Precision matrix (Column Sum=1)





Logistic Regression With
Class balancing Hyper
parameter tuning

```
[ ]: #Logistic Regression #With
      Class balancing #Hyper
      parameter tuning

alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2',
    ↪ loss='log', random_state=42)
    clf.fit(train_df, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_df, y_train)
    sig_clf_probs = sig_clf.predict_proba(cv_df)
    cv_log_error_array.append(log_loss(y_cv, sig_clf_probs, labels=clf.
    ↪ classes, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use
    ↪ log-probability estimates
    print("Log Loss :", log_loss(y_cv, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
```

```


for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
    penalty='l2', loss='log', random_state=42)
clf.fit(train_df, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df, y_train)

predict_y = sig_clf.predict_proba(train_df)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
    ", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_df)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
    log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_df)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
    ", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

for alpha = 1e-06
Log Loss : 0.19113171780419833
for alpha = 1e-05
Log Loss : 0.18584508186731702
for alpha = 0.0001
Log Loss : 0.16403903917983137
for alpha = 0.001
Log Loss : 0.1559052882112476
for alpha = 0.01
Log Loss : 0.15610417279373368
for alpha = 0.1
Log Loss : 0.161110735323708
for alpha = 1
Log Loss : 0.17351882883436703
for alpha = 10
Log Loss : 0.18729152018312847
for alpha = 100
Log Loss : 0.19001110697968235

```

For values of best alpha = 0.001 The train log loss is: 0.1727204896809822

For values of best alpha = 0.001 The cross validation log loss is:
0.1559052882112476


For values of best alpha = 0.001 The test log loss is: 0.1661671384860418

```
[ ]: clf = SGDClassifier(class_weight=balanced, alpha=alpha[best_alpha],
    ↪penalty='l2', loss='log', random_state=42)
    predict_and_plot_confusion_matrix(train_df.values, y_train.values, cv_df.
    ↪values, y_cv.values, clf)
```


Log loss : 0.1559052882112476

Number of mis-classified points : 163.63090128755366


..... Confusion matrix



Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)



Random Forest Classifier

Hyper paramter tuning

```
[ ]: import pickle
pickle.dump(sig_clf, open('/content/drive/MyDrive/final_prediction.pickle', 'wb'))
pickle.dump(scaler, open('/content/drive/MyDrive/scaler.pickle', 'wb'))

alpha = [100, 200, 500, 1000, 2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i, "and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini',
        max_depth=j, random_state=42, n_jobs=-1)
        clf.fit(train_df, y_train)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_df, y_train)
        sig_clf_probs = sig_clf.predict_proba(cv_df)
        cv_log_error_array.append(log_loss(y_cv, sig_clf_probs, labels=clf.
        classes_, eps=1e-15))
        print("Log Loss :", log_loss(y_cv, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[:, None], np.array(max_depth) [None]).ravel()
ax.plot(features, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
```

```

ax.annotate((alpha[int(i/2)], max_depth[int(i/2)], str(txt)), _
            (features[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], _
                             criterion='gini', max_depth=max_depth[int(best_alpha/2)], random_state=42, _
                             n_jobs=-1)
clf.fit(train_df, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_df, y_train)

predict_y = sig_clf.predict_proba(train_df)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train _
      log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_df)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross _
      validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, _
      eps=1e-15))
predict_y = sig_clf.predict_proba(test_df)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test _
      log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

for n_estimators = 100 and max depth = 5
Log Loss : 0.17485643131115827
for n_estimators = 100 and max depth = 10
Log Loss : 0.17689845098260562
for n_estimators = 200 and max depth = 5
Log Loss : 0.1749253413215618
for n_estimators = 200 and max depth = 10
Log Loss : 0.17412827734558625
for n_estimators = 500 and max depth = 5
Log Loss : 0.1679919273674284
for n_estimators = 500 and max depth = 10
Log Loss : 0.17209225434292702
for n_estimators = 1000 and max depth = 5
Log Loss : 0.16512473923084903
for n_estimators = 1000 and max depth = 10
Log Loss : 0.17073808391824918
for n_estimators = 2000 and max depth = 5
Log Loss : 0.16457364744805636
for n_estimators = 2000 and max depth = 10

```

Log Loss : 0.17070997401001095

For values of best estimator = 2000 The train log loss is: 0.15018613000009057

For values of best estimator = 2000 The cross validation log loss is:
0.16457364744805636

For values of best estimator = 2000 The test log loss is: 0.14752391650359561


Testing model with best hyper parameters

```
[ ]: clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)],  
                                criterion='gini', max_depth=max_depth[int(best_alpha/2)], random_state=42,  
                                n_jobs=-1)  
# predict_and_plot_confusion_matrix(train_x_onehotCoding,  
train_y_cv_x_onehotCoding_cv_y, clf)  
predict_and_plot_confusion_matrix(train_df.values, y_train.values, cv_df.  
values, y_cv.values, clf)
```


Log Loss : 0.17070997401001095

Number of mis-classified points : 163.42060085836908


..... Confusion matrix



..... Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)



##Conclusions: 1. Among 3 ML Models , Random Forest is the best ML Model for our task.
 2. Train and test performance of Model can be further improved by using Deep Learning models ,However at the Cost of computational expense.

```
[ ]: import pickle
pickle.dump(sig_clf, open('/content/drive/MyDrive/final_prediction.pickle', 'wb'))
pickle.dump(scaler, open('/content/drive/MyDrive/scaler.pickle', 'wb'))
```

```
[ ]:
```

```
[ ]:
```